



Using OVP Models in SystemC TLM2.0 Platforms

Imperas Software Limited

Imperas Buildings, North Weston,
Thame, Oxfordshire, OX9 2HA, UK
docs@imperas.com



Author:	Imperas Software Limited
Version:	2.3
Filename:	OVPsim_Using_OVP_Models_in_SystemC_TLM2.0_Platforms.doc
Project:	Using OVP Models in SystemC TLM2.0 Platforms
Last Saved:	Thursday, 25 November 2021
Keywords:	

Copyright Notice

Copyright © 2021 Imperas Software Limited All rights reserved. This software and documentation contain information that is the property of Imperas Software Limited. The software and documentation are furnished under a license agreement and may be used or copied only in accordance with the terms of the license agreement. No part of the software and documentation may be reproduced, transmitted, or translated, in any form or by any means, electronic, mechanical, manual, optical, or otherwise, without prior written permission of Imperas Software Limited, or as expressly provided by the license agreement.

Right to Copy Documentation

The license agreement with Imperas permits licensee to make copies of the documentation for its internal use only. Each copy shall include all copyrights, trademarks, service marks, and proprietary rights notices, if any.

Destination Control Statement

All technical data contained in this publication is subject to the export control laws of the United States of America. Disclosure to nationals of other countries contrary to United States law is prohibited. It is the reader's responsibility to determine the applicable regulations and to comply with them.

Disclaimer

IMPERAS SOFTWARE LIMITED, AND ITS LICENSORS MAKE NO WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, WITH REGARD TO THIS MATERIAL, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE.

Table of Contents

1	Introduction.....	5
1.1	Why use CpuManager or OVPSim?.....	5
1.2	Restrictions	5
1.3	Compiling Examples Described in this Document.....	5
2	How CpuManager works with SystemC TLM2.0	7
2.1	Platform construction.....	7
2.1.1	Naming.....	7
2.2	Processor models	7
2.2.1	Instructions/Sec and Quantum size.....	8
2.2.2	Guidelines for setting quantum and MIPS.....	8
2.2.2.1	Quantum vs Transaction delays.....	8
2.2.2.2	Other factors demanding a smaller quantum	8
2.2.2.3	Factors demanding a larger quantum.....	9
2.2.3	SystemC Stack Size	9
2.2.4	Direct Memory Interface Memory Access	9
2.2.5	Simulation artifacts	10
2.2.6	Delays in bus transactions.....	10
2.3	Peripheral models.....	10
2.3.1	Delays in bus transactions.....	10
2.4	Automatic generation of the TLM interface.	10
3	OVP OP header and source files.....	11
3.1	Building and Linking C++ Interfaces	12
3.1.1	Build Local TLM and OP++ Archive by Default.....	12
3.1.2	Optionally Use Global TLM Archive	13
3.1.3	Manually Building the Global Archives	13
4	Generating API TLM Interface files.....	15
4.1	Peripheral Models	15
4.2	Processor Models	15
5	Example Platform	17
5.1	Compilation.....	19
5.2	Building an application	21
5.3	Running a platform	21
5.4	Platform Construction Options	22
5.4.1	Processor Options	22
5.4.1.1	Setting a variant	22
5.4.1.2	Instruction Tracing.....	22
5.4.1.3	Application debug.....	22
5.4.1.4	Setting the simulation time slice (quantum)	23
5.4.1.5	Simulated Exceptions.....	23
5.4.1.6	Loading intercept libraries	24
5.4.2	Peripheral Options	24
5.4.2.1	Peripheral diagnostics.	24
5.4.3	DMI.....	24
6	Deviations from TLM2.0 LRM	26
6.1	Data Endian in TLM transactions	26

6.2	Modeling of Interrupts	26
7	Tracing TLM activity.....	27
8	What can go wrong	28
8.1	Spaces in filenames.....	28
8.2	OVPsim version incompatibilities	28
8.3	Environment problems.....	28
8.4	Compiling OSCI SystemC 2.2.0 with later versions of gcc	28

1 Introduction

This document describes the use of OVP models in systemC TLM2.0 simulation platforms using the OP C++ API. Earlier releases of the Imperas simulator used the deprecated (still supported) ICM C++ API.

CpuManager and OVPSim are dynamic linked libraries (.so suffix on Linux, .dll suffix on Windows) implementing Imperas simulation technology. The shared objects contain implementations of the OP interface functions described in "OVPSim and CpuManager User Guide". The OP functions enable instantiation, interconnection and simulation of complex multiprocessor platforms containing arbitrary shared memory topologies and peripheral devices.

CpuManager is one of the commercial products available from Imperas. OVPSim available from www.ovpworld.org

It is assumed that you are familiar with C++, SystemC and TLM2.0 technology.

Please refer to the "OVPSim and CpuManager User Guide" for more details of OVPSim and CpuManager products.

1.1 Why use CpuManager or OVPSim?

OVPSim and CpuManager (hereafter referred to as just CpuManager) have access to a rich source of fast, qualified processor models and to a constantly growing list of peripheral models. Using CpuManager in your SystemC TLM2.0 simulation gives access to these high performance models, and to associated software development tools with very little extra effort.

1.2 Restrictions

CpuManager is a high speed instruction-accurate processor and platform simulator. It is not intended for cycle-accurate or pin-level simulation. For this reason the TLM2.0 interface uses the TLM2.0 "loosely timed" (LT) model. Attempting to use other models will give incorrect results.

CpuManager allows the free-running of each processor for a large number of instructions rather than advancing all processors in lock-step. If your simulation uses TLM2.0 models which rely on lock-step operation you will need to reduce to one the number of instructions which are run in each step.

1.3 Compiling Examples Described in this Document

This documentation is supported by C++ code samples in an `Examples` directory, available as part of an OVPSim installation, by download from the www.ovpworld.org website or as part of an Imperas installation.

The example uses the OR1K processor model and tool chain. The model

is included as part of the OVPsim or Imperas installation. The toolchain is available by free download from the www.ovpworld.org website.

For Windows environments, Imperas recommends using MSYS2 (www.msys2.org).

SystemC TLM2.0 models can be used on Windows with MinGW/MSys (since SystemC release v2.3.0). It is assumed that users of this environment will be familiar with C++, SystemC, TLM2.0 and will have obtained this software from <http://accellera.org/downloads/standards/systemc> or similar.

All version of SystemC TLM2.0 since v2.3.0 have been used by Imperas, the latest version supported is SystemC release v2.3.3.

2 How CpuManager works with SystemC TLM2.0

2.1 Platform construction

An OVP model is provided with a TLM2.0 interface in the form of C++ source and header files. The headers should be included in the TLM2.0 platform source. They define a SystemC classes for processor, peripheral, MMC and other models. These classes are instanced and configured to make the specific models required in the platform.

An archive of the compiled classes is provided but this might not be compatible with your C++ compiler, so you should be prepared to compile them for use in your platform.

The OVP SystemC interface uses the following model classes:

Model	Class	File
Root Module	tlmModule	tlmModule.hpp
Processor	tlmProcessor	tlmProcessor.hpp
Peripheral	tlmPeripheral	tlmPeripheral.hpp
MMC	tlmMMC	tlmMmc.hpp

Additional interface classes are defined :

Function	Class	File
Bus master interface	tlmBusMasterPort	tlmBusPort.hpp
Bus slave interface	tlmBusSlavePort	tlmBusPort.hpp
Dynamic slave interface	tlmBusDynamicSlavePort	tlmBusPort.hpp
Net input	tlmNetInputPort	tlmNetPort.hpp
Net output	tlmNetOutputPort	tlmNetPort.hpp
Time synchronization	time_advance	tlmModule.hpp

The root module is required by OVPsim to be the parent of each model. The SystemC integration synchronizes the SystemC scheduler with the OVP scheduler so that peripheral models that use time-related functions have the correct relationship to SystemC's clock.

2.1.1 Naming

The TLM2.0 interface uses the SystemC method `sc_object::name()` to create a dot-separated hierarchical name for each CpuManager model instance.

2.2 Processor models

An OVP processor model has one or more execution units, defined as having it's own program counter capable of executing a program independent of other execution units. Each execution unit is run in a SystemC thread. The thread executes a calculated number of instructions on the processor without advancing SystemC time. Each instruction may or may not cause TLM2.0 transactions to be propagated to other components in the platform. A transaction may call SystemC `wait()`. When the allotted instructions have

completed, the thread calls `SystemC wait()` to advance time. The processor threads are executed in an order determined by the SystemC scheduler.

2.2.1 Instructions/Sec and Quantum size

To use OVP models, SystemC must instantiate one `tlmModule` object. This object keeps the quantum period which sets how long each processor model instance waits before running again.

Each processor model instance keeps a figure which controls the effective number of instructions per second (IPS) executed by the model. It uses this and the quantum period to decide how many instructions to run in each quantum.

The default quantum period is 1mS. The default IPS is 100,000,000. Thus, by default, a processor runs 100,000 instructions per quantum (this matches OVPsim's internal scheduler used in a non-SystemC environment).

To change the quantum period use the `tlmModule.quantum()` method in your platform constructor. The effective frequency of a processor instance is set by the `mips` parameter which is passed to the constructor of the `tlmProcessor` class.

2.2.2 Guidelines for setting quantum and MIPS

A processor's MIPS should be set to give the correct clock frequency with respect to other models in the simulation. Note that if no other models accurately represent time, then setting the IPS will not affect the behavior of the simulation, merely the reported statistics.

Setting the quantum period is a compromise; a smaller quantum yields a more accurate representation of reality, a larger quantum achieves higher simulation speed.

2.2.2.1 Quantum vs Transaction delays

As described above, by default, the processor model interface runs a larger number of instructions that are assumed to be un-timed, then executes a delay to give a crude model of clock frequency. If the TLM target models add a per-transaction delay, then the user should change the interface to run a much smaller number of instructions per quantum, possibly one. Also, the per-instruction delay might need to be reduced to take account of delays in the target models.

2.2.2.2 Other factors demanding a smaller quantum

To avoid gross functional errors, the quantum period for a processor must be shorter than the shortest time delay modeled in any peripheral device with which the processor interacts.

Similarly, if two processors are communicating using shared memory, the number of instructions per quantum should be less than the number of instructions taken to make each communication.

2.2.2.3 Factors demanding a larger quantum

A short quantum results in poor simulation performance. However, this is only of concern if you intend to simulate many instructions. As a guideline, the SystemC scheduler takes at best a few hundred instructions to start a processor's quantum, so as the instructions per quantum is reduced to this number, the performance will be dominated by the scheduler and not the model.

Scenario	Dominating factor	Quantum
Booting Linux. Functional (not cycle accurate) peripheral models.	Simulation speed of processor model	$\geq 1\text{mS}$
Programming a graphics controller. Cycle-accurate GPU.	Simulation speed of GPU	$< 1\mu\text{S}$
Developing a UART driver. Uart has 1mS character rate.	Accuracy of interaction	$< 1\text{mS}$

2.2.3 SystemC Stack Size

The `tlmProcessor` model requires an increased thread stack depth. The function `set_stack_size()` is used to override the default SystemC thread stack size.

2.2.4 Direct Memory Interface Memory Access

TLM2.0 allows comprehensive modeling of bus transactions, but each transaction takes significant simulation time. Direct Memory Interface (DMI) allows negotiation between two TLM2.0 models so that an initiator can directly access target memory, bypassing the TLM2.0 mechanism. This completely sacrifices timing accuracy for simulation speed.

The TLM2.0 OVP interface uses DMI negotiation by default. In practice, a processor with a fixed program memory will execute one code-fetch via TLM2.0. The DMI hint in that transaction will allow `CpuManager` to map the program memory into the processor's address space so that subsequent code-fetches do not use TLM2.0 transactions. DMI can be enabled or disabled by calling the `dm_i()` method on a processor or memory instance. See section 5.4.3.

Note that a simulator that uses code translation must necessarily cache translated code in the simulator. If the original code memory is subsequently modified by a mechanism outside the simulator, the simulator must be notified so that the code can be re-translated. `CpuManager` supports the standard DMI invalidation mechanism.

When DMI is turned off on a processor, cached DMI regions are removed for its entire code and data address spaces

When DMI is turned off on a memory, the cached DMI region is removed from connected processors for the region that addresses the memory.

When DMI is turned on again, the models will re-negotiate DMI as each transaction occurs.

2.2.5 Simulation artifacts

CpuManager performs dynamic code translation for simulation efficiency. A processor will therefore pre-fetch each code location (up to the next jump or branch) once before it begins executing code. The pre-fetches use the TLM `transport_dbg()` method rather than the regular `transport()` method, to distinguish artifact accesses from the real accesses.

For this reason, bus and memory models **must support the `transport_dbg()` method**. The simplest way to do this is to give the `transport()` and `transport_dbg()` methods the same behavior. However if a model counts or otherwise reports bus traffic, it should not do so in the `transport_dbg()` method.

2.2.6 Delays in bus transactions

The LT (loosely timed) TLM model allows bus transactions to take time. It is legal for a CpuManager processor model to be stalled by a bus transaction that takes time – some part of a bus target calls SystemC wait, which stops the processor model in its bus access. Other scheduled tasks, including other processor models can run while this is happening. If a processor model is used in this way, it is not appropriate to use the supplied TLM processor wrappers – the user will need to write a wrapper that schedules the processor one instruction at a time.

2.3 Peripheral models

During simulation, peripheral models can be activated in three ways:

- TLM2.0 transaction. A TLM2.0 transaction by another model is propagated to this model which results in a bus read or write.
- Elapsed time. Each time SystemC advances time, it notifies CpuManager, which will activate any peripheral waiting for that time to occur.
- Net propagation. A write to a net (implemented using the SystemC `analysis port`) will be propagated by SystemC to any peripheral models connected.

CpuManager also requests notification at the beginning and end of simulation to trigger OVP peripheral model BHM simulation-start and simulation-end special events.

2.3.1 Delays in bus transactions

The LT (loosely timed) TLM model allows bus transactions to take time. It is NOT legal for a CpuManager peripheral model to be stalled by a bus transaction that takes time – CpuManager expects peripheral model code to run in “zero time” – scheduling another model while a peripheral model is stalled will lead to unpredictable behaviour and memory corruption.

2.4 Automatic generation of the TLM interface.

The Imperas Model Generator (igen) can be used to generate code to build a SystemC platform (one that uses only OVP models).

Refer to the Imperas_Model_Generator_Guide.doc.

3 OVP OP header and source files

The OP API, used by both CpuManager and OVPSim, is defined by several header files within the Imperas tool release tree or download from www.ovpworld.org :

Common Definitions

Standard types ImpPublic/include/host/impTypes.h

OP API Definitions

C API functions ImpPublic/include/host/op/op.h

C++ API functions ImpPublic/include/host/op/op.hpp
 ImpPublic/source/host/op/op.cpp

OP API Link library

C API functions bin/<host architecture>/libRuntimeLoader.so
 bin/<host architecture>/libRuntimeLoader.dll

C++ API functions¹ bin/<host architecture>/libRuntimeLoader++.so
 bin/<host architecture>/libRuntimeLoader++.dll

Required by a module: **tlmModule class**

ImpPublic/source/host/tlm/tlmModule.cpp
ImpPublic/include/host/tlm/tlmModule.hpp

Required by a processor interface: **tlmProcessor class**

ImpPublic/source/host/tlm/tlmProcessor.cpp
ImpPublic/include/host/tlm/tlmProcessor.hpp

Required by a peripheral interface: **tlmPeripheral class**

ImpPublic/source/host/tlm/tlmPeripheral.cpp
ImpPublic/include/host/tlm/tlmPeripheral.hpp

Required by TLM platforms written by **igen**:

Generic Memory Model

ImpPublic/source/host/tlm/tlmDenseMemory.cpp (dense)
ImpPublic/include/host/tlm/tlmDenseMemory.cpp (dense)
ImpPublic/source/host/tlm/tlmMemory.hpp (sparse)
ImpPublic/include host/tlm/tlmMemory.hpp (sparse)

Generic Bus Decoder

ImpPublic/source/host/tlm/tlmDecoder.cpp
ImpPublic/include/host/tlm/tlmDecoder.hpp

¹ These are obsolete, but provided for backwards compatibility; by default, ImpPublic/source/host/op.op.cpp is now built using the local host C++ compiler by the Imperas-provided Makefiles. See section **Error! Reference source not found. Error! Reference source not found.** for details.

3.1 Building and Linking C++ Interfaces

When building a SystemC or SystemC TLM platform a C++ compiler/linker is used. The version used must be consistent across all libraries and executables built.

The TLM interface and the OP C++ interface source files are provided so they may be built locally, thus ensuring they are consistent with the SystemC library they are linked with.

By default, these files are automatically compiled and put into archive libraries local to the platform being built when the Imperas-provided Makefiles are used.

To use a single global library for all platforms just set the environment variables `IMPERAS_TLM_SUPPORT_ARCHIVE` and `IMPERAS_OP_CPP_LIBRARY` respectively to the full path of the TLM and OP C++ archives to use. The archives will automatically be built if they do not already exist.

The following sections show examples of this. (They use the example provided in *IMPERAS_HOME/Examples/PlatformConstruction/SystemC_TLM*)

3.1.1 Build Local TLM and OP++ Archive by Default

By default a local copy is built and linked:

```
$ cp -r IMPERAS_HOME/Examples/PlatformConstruction/SystemC_TLM .
$ cd SystemC_TLM
$ make -C platform_cpp VERBOSE=1
make: Entering directory './SystemC_TLM/platform_cpp'
# INFO: Using local TLM support archive tlmSupport/tlm.a; Set IMPERAS_TLM_SUPPORT_ARCHIVE
to specify an alternative library built from source files in
<IMPERAS_HOME>/ImpPublic/source/host/tlm"
# iGen Create OP TLM2.0 EXAMPLE PLATFORM platform.cpp
...
# OP++: Create local OP CPP library directory opLib/
mkdir -p opLib/
...
# OP++: Building Support Archive libOP++.a
...
# Copy TLM Support files from <IMPERAS_HOME>/ImpPublic/source/host/tlm
...
# Build TLM Archive tlmSupport/tlm.a
...
# TLM Linking platform.Linux64.exe
mkdir -p .
g++ -o platform.Linux64.exe Build/Linux64/usr/platform.o tlmSupport/tlm.a
/home/nda/SystemC/install/systemc-2.3.3/lib-linux64/libsystemc.a -
L<IMPERAS_HOME>/bin/Linux64 -lRuntimeLoader -lisl -lm64 -fPIC -LopLib/ -lOP++ -lpthread
...
```

This should be sufficient for most users.

3.1.2 Optionally Use Global TLM Archive

If many different executables are built in your environment you may wish to have only a single copy of these archives present.

By setting the environment variables specifying a path for the tlm.a and/or libOP++.a archives we can re-use the same archives with multiple executables, and avoid the overhead of re-building them for each executable:

```
$ export IMPERAS_TLM_SUPPORT_ARCHIVE=~/.tlmSupport/tlm.a
$ export IMPERAS_OP_CPP_LIBRARY=~/.opSupport/libOP++.a
$ cp -r IMPERAS_HOME/Examples/PlatformConstruction/SystemC_TLM .
$ cd SystemC_TLM
$ make -C platform_cpp VERBOSE=1
...
# TLM Linking platform.Linux64.exe
mkdir -p .
g++ -o platform.Linux64.exe Build/Linux64/usr/platform.o ~/.tlmSupport/tlm.a
/home/nda/SystemC/install/systemc-2.3.3/lib-linux64/libsystemc.a -
L<IMPERAS_HOME>/bin/Linux64 -lRuntimeLoader -lisl -m64 -fPIC -L~/.opSupport/ -lOP++ -
lpthread
...
```

Note that the global libraries indicated by the environment variables will be automatically built if they do not already exist, when using the provided Makefiles.

3.1.3 Manually Building the Global Archives

To manually build the archives yourself, simply make a local copy of the source directories from the installation and run make in them:

```
$ mkdir tlmSupport
$ cp IMPERAS_HOME/ImpPublic/source/host/tlm/* tlmSupport
$ make -C tlmSupport
make: Entering directory '~/tmp/SystemC_TLM/tlmSupport'
# TLM Support Compiling Build/Linux64/usr/tlmBusDynamicSlavePort.o
# TLM Support Compiling Build/Linux64/usr/tlmBusMasterPort.o
# TLM Support Compiling Build/Linux64/usr/tlmBusSlavePort.o
# TLM Support Compiling Build/Linux64/usr/tlmDMISlave.o
# TLM Support Compiling Build/Linux64/usr/tlmDecoder.o
# TLM Support Compiling Build/Linux64/usr/tlmDenseMemory.o
# TLM Support Compiling Build/Linux64/usr/tlmMemory.o
# TLM Support Compiling Build/Linux64/usr/tlmMmc.o
# TLM Support Compiling Build/Linux64/usr/tlmModule.o
# TLM Support Compiling Build/Linux64/usr/tlmNetInputPort.o
# TLM Support Compiling Build/Linux64/usr/tlmPeripheral.o
# TLM Support Compiling Build/Linux64/usr/tlmProcessor.o
# TLM Support Archive tlm.a
ar: creating tlm.a
# TLM Support NM tlm.nm.out
# TLM Support Library tlm.a
make: Leaving directory '~/tmp/SystemC_TLM/tlmSupport'
```

```
$ mkdir opSupport
$ cp IMPERAS_HOME/ImpPublic/source/host/op/* opSupport
$ make -C opSupport
make: Entering directory '~/opSupport'
# OP++: Depending Build/Linux64/usr/op.d
# OP++: Compiling Build/Linux64/usr/op.o
# OP++: Building Support Archive libOP++.a
```

```
ar: creating libOP++.a  
# OP++: Created C++ Support archive library: libOP++.a  
make: Leaving directory '~/opSupport'
```

4 Generating API TLM Interface files

The VLNV library provides OP API TLM2.0 interface files for the default model configurations.

The next two sections show how to accomplish this.

4.1 Peripheral Models

Each peripheral configuration requires a separate TLM interface file to match the specific configuration of, for example, interrupt connections.

The following shows the generation of the **default** TLM2.0 interface file for the RISC-V peripheral PLIC.

```
$ igen.exe \  
  -modelvender riscv.ovpworld.org -modellibrary peripheral -modelname PLIC \  
  -writetlm plic.igen.hpp
```

This peripheral can be configured to support different numbers of sources and targets

```
$ igen.exe \  
  -modelvender riscv.ovpworld.org -modellibrary peripheral -modelname PLIC \  
  -writetlm plic_s2_t4.igen.hpp \  
  -setparameter num_sources=2 \  
  -setparameter num_targets=4
```

4.2 Processor Models

Each processor configuration requires a separate TLM interface file to match the specific configuration of, for example, bus ports and interrupt connections.

The following shows the generation of the **default** TLM2.0 interface file for the ARM processor configured for variant ARMv5TEJ.

```
$ igen.exe \  
  -modelvender arm.ovpworld.org -modellibrary processor -modelname arm \  
  -variant ARMv5TEJ -writetlm arm_ARMv5TEJ.igen.hpp \  
  -userheader ImperasLib/fileheaders/refArmApache.txt
```

Additional configuration parameters can be assigned by adding *-setparameter* arguments.

For example, generate the TLM2 interface for the ARM CortexAMPx2 processor with additional interrupt lines.

```
$ igen.exe \  
-modelvendor arm.ovpworld.org -modellibrary processor -modelname arm \  
-variant CortexA9MPx2 \  
-setparameter override_GICD_TYPER_ITLines=6\  
-writetlm arm_Cortex-A9MPx2-GICD_TYPER_ITLine-6.igen.hpp \  
-userheader ImperasLib/fileheaders/refArmApache.txt
```


5 Example Platform

An example TLM2.0 platform is provided in
`Examples/PlatformConstruction/SystemC_TLM`

Figure 1: Example TLM2.0 Platform Block Diagram

The source files for this example are in
`Examples/PlatformConstruction/SystemC_TLM/platform_cpp`

The complete example platform has all memory mapped to TLM memory. At runtime this will be accessed using DMI, if possible.

The platform-specific code is in the included file `platform.sc_constructor.cpp`:

The model header files are included in the platform header:

```
#include "tlm/tlmModule.hpp"
#include "tlm/tlmDecoder.hpp"
#include "tlm/tlmMemory.hpp"
#include "tlm/tlmProcessor.hpp"
#include "tlm/tlmPeripheral.hpp"
```

In the platform class, platform components are instantiated (template parameters are supplied to template components).

```
class simple : public sc_core::sc_module {
public:
    simple (sc_core::sc_module_name name);

    tlmModule          Platform;
    tlmDecoder          bus1;
    tlmRam              ram1;
    tlmRam              ram2;
    tlmProcessor        cpul;
```

```

    tlmPeripheral      uart1;

    const char *pathForcpul() {
        return opVLNVString(0, "ovpworld.org", "processor", "orlk", "1.0", OP_PROCESSOR, 1);
    }

    const char *pathForuart1() {
        return opVLNVString(0, "national.ovpworld.org", "peripheral", "16450", "1.0", OP_PERIPHERAL, 1);
    }

    params paramsForuart1() {
        params p;
        p.set("outfile", "uart1.log");
        return p;
    }

    params paramsForcpul() {
        params p;
        p.set("defaultsemihost", true);
        return p;
    }
}; /* simple */

```

The `pathForcpul()` and `pathForuart1()` functions provide the path to the models from the VLNV.

The `paramsForuart1()` function specifies configuration parameters for the uart model, in this case defining a file in which the uart's output is to be written.

The `paramsForcpul()` function specifies configuration parameters for the processor model, in this case turning on the default semihost library to intercept system calls, sending text output to the simulator's output stream.

Calls to the sub-constructors are put before the body of the constructor.

```

simple::simple ( sc_core::sc_module_name name)
: sc_module (name, paramsForRoot())
, Platform ("")
, bus1 (Platform, "bus1", 2, 3)
, ram1 (Platform, "ram1", 0xfffff)
, ram2 (Platform, "ram2", 0xffff)
, cpul (Platform, pathForcpul(), "cpul", paramsForcpul())
, uart1(Platform, pathForuart1(), "uart1", paramsForuart1())
{

```

Connections (binding) between components occur in the body of the constructor.

```

    // bus1 masters
    bus1.connect(cpul, "INSTRUCTION", 32);
    bus1.connect(cpul, "DATA", 32);

    // bus1 slaves
    uart1.connect(bus1, "bport1", 0x90000000, 0x90000007); // Peripheral (0)

    bus1.connect(mem1.sp1, 0x0, 0xffffffff); // Memory (1)
    bus1.connect(mem2.sp1, 0x90000008, 0xffffffff); // Memory (2)

    // Net connections
    uart1.connect("intOut", cpul, "intr2");

```

The file `platform.cpp` defines `sc_main` which instances the simulator session, the standard command line parser and the platform class `simple`:

```
...
int sc_main (int argc, char *argv[] ) {
    session s;
    parser p(argc, (const char**) argv);

    simple simple("simple");

    sc_start();
    return 0;
}
```

The TLM interface `tlmProcessor` is derived from the OP C++ `processor` class so that all methods of this class can be applied to the TLM processor instantiation. The `tlmPeripheral` is derived from OP C++ `peripheral` so you can similarly use its methods.

The SystemC library provides `main()` that calls `sc_main()` after the constructors and starts the SystemC scheduler.

5.1 Compilation

The above example² was compiled under Windows using MinGW/MSys³ and on Linux using GCC.

To build the example, follow these steps:

On Windows

- Obtain and install MinGW/MSys

On Windows and Linux

- Obtain and install SystemC v2.3 source or above that support MinGW build
- Obtain and install the OR1K tool-chain (this example uses an or1k processor model)
- Obtain and install OVPsim or the Imperas professional tools and configure the Imperas environment as described in the Installation guide

All compilation is performed in the Linux or MinGW/MSys command shell

```
> cd <temp directory>
> cp -r $IMPERAS_HOME/Examples/PlatformConstruction/SystemC_TLM .
> cd SystemC_TLM/platform_cpp
```

Specify the locations of your SystemC and TLM2.0 releases. For Example

```
> export SYSTEMC_HOME=C:/SystemC/systemc-2.3.3
```

Compile the example platform and interfaces:

The Makefile includes standard Makefiles⁴ that provide rules and targets.

² The old ICM API example `IMPERAS_HOME/Examples/Platforms/ICM/SystemC_TLM2.0` includes the batch file `platform_cpp/compile.msvc.bat` that uses `nmake` in an MSVC command prompt to build the platform. This is no longer supported with the usage of MSYS/MinGW.

³ Since SystemC release v2.3.0 support to build with MinGW on Windows has been included. At this time OVP moved all SystemC TLM2 examples and demos from building with MSVC to building with MinGW.

```
# Common variables
include $(IMPERAS_HOME)/ImperasLib/buildutils/Makefile.common
# Target and Rules C to executable
include $(IMPERAS_HOME)/ImperasLib/buildutils/Makefile.TLM.platform
# Target and Rules iGen to C
include $(IMPERAS_HOME)/ImperasLib/buildutils/Makefile.TLM.igen
```

Makefile.common : Defines common variables
Makefile.TLM.platform : Provides main target and rules to build executable from C source files
Makefile.TLM.igen : Provides target and rules to generate C source from iGen input files

To build:

```
> make
# iGen Create OP TLM2.0 EXAMPLE PLATFORM platform.cpp
# TLM Depending Build/Linux32/usr/platform.d
# TLM Compiling Build/Linux32/usr/platform.o
# TLM Linking platform.Linux32.exe
```

To see the commands that are being performed, use **VERBOSE=1**

```
> make realclean
> make VERBOSE=1
# iGen Create OP TLM2.0 EXAMPLE PLATFORM platform.cpp
igen.exe --quiet --nobanner --excludem GPT_NH --excludem GPT_UFNR \
  --op \
  --batch platform.tlm.tcl \
  --writetlm platform.cpp \
  --userheader <IMPERAS_HOME>/ImperasLib/fileheaders/refESLA.txt \
  --overwrite
# TLM Depending Build/Linux32/usr/platform.d
mkdir -p Build/Linux32/usr
g++ -MM platform.cpp -I<IMPERAS_HOME>/ImpPublic/include/host -I<SYSTEMC_HOME>/include -
I<IMPERAS_HOME>/ImperasLib/source -ggdb -Wno-long-long -Wall -Werror -
DSC_INCLUDE_DYNAMIC_PROCESSES -D_CRT_SECURE_NO_WARNINGS -D_CRT_SECURE_NO_DEPRECATED -
I<IMPERAS_HOME>/ImpPublic/include/host -I<IMPERAS_HOME>/ImpProprietary/include/host -Wall
-Werror -O0 -g -gdwarf-2 -m32 -D_GLIBCXX_USE_CXX11_ABI=0 -MT
Build/Linux32/usr/platform.o -MF Build/Linux32/usr/platform.d
# TLM Compiling Build/Linux32/usr/platform.o
mkdir -p Build/Linux32/usr
g++ -c -o Build/Linux32/usr/platform.o platform.cpp -
I<IMPERAS_HOME>/ImpPublic/include/host -I<SYSTEMC_HOME>/include -
I<IMPERAS_HOME>/ImperasLib/source -ggdb -Wno-long-long -Wall -Werror -
DSC_INCLUDE_DYNAMIC_PROCESSES -D_CRT_SECURE_NO_WARNINGS -D_CRT_SECURE_NO_DEPRECATED -
I<IMPERAS_HOME>/ImpPublic/include/host -I<IMPERAS_HOME>/ImpProprietary/include/host -Wall
-Werror -O0 -g -gdwarf-2 -m32 -D_GLIBCXX_USE_CXX11_ABI=0
# TLM Linking platform.Linux32.exe
mkdir -p .
g++ -o platform.Linux32.exe Build/Linux32/usr/platform.o -L<IMPERAS_HOME>/bin/Linux32 -
lRuntimeLoader -lisl -m32 <SYSTEMC_HOME>/lib-linux/libsystemc.a -
L<IMPERAS_HOME>/bin/Linux32 -lRuntimeLoader -lisl -m32 -lRuntimeLoader++ -lpthread
test -f /usr/bin/execstack; \
  if [ $? -eq 0 ]; then \
    execstack --clear-execstack platform.Linux32.exe; \
  fi
```

⁴ The order of inclusion is important

5.2 Building an application

A toolchain to allow an application to be cross compiled for the OR1K processor can be obtained from www.ovpworld.org.

The application code cross compilation is supported for a MINGW shell on Windows. Refer to the document “OVPsim_Installation_and_Getting_Started” for installation and use information.

To build the application in a Linux or in an MSYS shell on Windows, follow these steps:

- Go to a copy of the application directory in the example
- Execute the provided Makefile

On Windows:

```
> cd ../application
> make
```

On Linux:

```
> cd ../application
> make
> cd ..
```

This will build the executable *int.OR1K.elf*

5.3 Running a platform

Run the platform, passing the application program:

```
platform_cpp/platform.$IMPERAS_ARCH.exe -program application/int.elf
```

The output from the run should be:

```
SystemC 2.3.3-Accellera --- Sep 30 2020 09:02:20
Copyright (c) 1996-2018 by all Contributors,
ALL RIGHTS RESERVED

TEST: main starts
TEST: Initialize:
TEST: Enable UART:
Interrupt Handler 0x02 (1)
Character sent
TEST: main send string
TEST: Send String: Hello World

Send char H (0x48)
Interrupt Handler 0x02 (2)
Character sent
Send char e (0x65)
Interrupt Handler 0x02 (3)
... etc ...
Character sent
Send char l (0x6c)
Interrupt Handler 0x02 (11)
Character sent
Send char d (0x64)
Interrupt Handler 0x02 (12)
Character sent
Send char
(0x0a)
```

```
Interrupt Handler 0x02 (13)
Character sent
TEST: main done
```

5.4 Platform Construction Options

OVP models can be instantiated in a TLM2.0 platform exactly like other models. However, there are many features in CpuManager which are available through the TLM2.0 interface. Some commonly used features are listed here. For details, refer to OVPsim_and_CpuManager_User_Guide.doc. The code examples are given in the context of the worked example platform.

5.4.1 Processor Options

5.4.1.1 Setting a variant

By default, a processor model will execute as an Instruction Set Architecture (ISA) model which represents the instructions but not necessarily the configuration of the processor. The `variant` parameter is used to select a different configuration; refer to the processor model specific documentation for the list of supported variants.

In the example, parameters can be set in the function `paramsForcpu1` (note that the `or1k` processor model does not have any variants):

```
params paramsForcpu1() {
    params p;
    p.set("defaultsemihost", true);
    p.set("variant", "variant1");
    return p;
}
```

5.4.1.2 Instruction Tracing.

Instruction tracing can be enabled from the simulator command line:

```
> platform_cpp/platform.$IMPERAS_ARCH.exe \
  -program application/int.elf
  -trace
```

Tracing can also be enabled by setting a parameter on the model:

```
params paramsForcpu1() {
    params p;
    p.set("defaultsemihost", true);
    p.set("variant", "variant1");
    p.set("trace", true);
    return p;
}
```

5.4.1.3 Application debug

The or1k processor shipped with this OVPsim release is provided with a (rather old) GNU gdb debugger. This can be started and connected to the or1k processor model from the simulator command line:

```
> platform_cpp/platform.$IMPERAS_ARCH.exe \
  -program application/int.elf \
  -gdbconsole
```

This will pop up a console containing a gdb already connected to the simulator.

Like all command line flags, the same effect can be achieved by setting a parameter, this time on the root module:

```
params paramsForRoot() {
    params p;
    p.set("gdbconsole", true);
    return p;
}

simple::simple ( sc_core::sc_module_name name)
: sc_module (name, paramsForRoot())
, Platform ("")
, bus1(Platform, "bus1", 2, 3)
, ram1 (Platform, "ram1", 0xfffff)
, ram2 (Platform, "ram2", 0xffff)
, cpul (Platform, "cpul", paramsForcpul())
, uart1(Platform, "uart1", paramsForuart1())
{
    ...
}
```

5.4.1.4 Setting the simulation time slice (quantum)

The quantum is set for the whole platform. All processor models use this value. See ImperasLib/source/ovpworld.org/modelSupport/tlmPlatform/1.0/tlm/tlmPlatform.hpp.

```
simple:: simple (sc_module_name name)
: Platform("simple", ...)
{
    Platform.setSimulationTimeSlice(0.0001);
}
```

5.4.1.5 Simulated Exceptions

By default an OVP processor model will notify the simulator if an exception (e.g. divide by zero or access alignment error) occurs. Set the `simulateexceptions` parameter to make the processor jump to its exception vector instead. This option is typically used when simulating an operating system. It is not used in a "bare metal" platform which has no code to handle the exception. See `op.h` and search for `OP_FP_`.

```
params paramsForcpul() {
    params p;
    p.set("simulateexceptions", true);
    return p;
}

simple::simple ( sc_core::sc_module_name name)
: sc_module (name, paramsForRoot())
, Platform ("")
```

```
, bus1( Platform, "bus1", 2, 3)
, ram1 (Platform, "ram1", 0xfffff)
, ram2 (Platform, "ram2", 0xffff)
, cpul (Platform, pathForcpul(), "cpul", paramsForcpul())
, uart1(Platform, pathForuart1(), "uart1", paramsForuart1())
{
    ...
}
```

5.4.1.6 Loading intercept libraries

We have already seen that this example uses the default intercept library for the or1k processor. Commented-out code can be included to load a library.

```
( in the class definition )
// extension          ext1;

( a private function in the platform class )
const char *getExt1Path(void) {
    return opVLNVString(
        0,
        "ovpworld.org",
        "semihosting",
        "or1kNewlib",
        "1.0",
        OP_EXTENSION,
        true
    );
}

( in the platform constructor )
,ext1 (cpul, getExt1Path(), "ext1")
```

5.4.2 Peripheral Options

5.4.2.1 Peripheral diagnostics.

Most peripherals are capable of producing diagnostic output with different levels of detail. Call the PSE method `diagnosticLevelSet ()`. See `bhm.h` : `bhmSetDiagnosticLevel()`.

```
uart1. diagnosticLevelSet (3);
```

5.4.3 DMI

As mentioned previously, OVP processor and memory models have DMI enabled by default. To turn DMI on or off use the `dmi` method on the memory or processor models, as illustrated in commented code in the platform constructor:

```
simple::simple (sc_module_name name)
: sc_module (name)
, Platform ("", paramsForRoot())
, bus1(Platform, "bus1", 2, 3)
, ram1 (Platform, "ram1", 0xfffff)
, ram2 (Platform, "ram2", 0xffff)
, cpul (Platform, pathForcpul(), "cpul", paramsForcpul())
, uart1(Platform, pathForuart1(), "uart1", paramsForuart1())
{
    bus1.connect(cpul, "INSTRUCTION", 32);
```



```
bus1.connect(cpu1.DATA);
bus1.connect(uart1, "bport1", 0x90000000, 0x90000007); // Peripheral (0)
bus1.connect(ram1.spl, 0x0, 0xffff); // Memory (1)
bus1.connect(ram2.spl, 0xffff0000, 0xffffffff); // Memory (2)
uart1.connect("intOut", cpu1, "intr0");

// By default DMI is turned on. Use these to turn it off
//ram1.dmi(0);
//ram2.dmi(0);
//cpu1.dmi(0);
}
```

6 Deviations from TLM2.0 LRM

6.1 *Data Endian in TLM transactions*

The contents of the data field in a TLM transaction is **target endian** rather than **host endian** as specified. This is due to an inconsistency in the TLM standard which makes efficient DMI otherwise hard to achieve.

6.2 *Modeling of Interrupts*

To model interrupt signals, the OVP interface to TLM2.0 uses the TLM analysis port rather than the SystemC net. The analysis port immediately propagates new values by function call, whereas the SystemC net requires the SystemC scheduler to cause propagation. As explained elsewhere in this document, for efficiency, the OVP processor model typically simulates thousands of instructions in one 'quantum' without the intervention of the SystemC scheduler. If these instructions change the value of an interrupt net, the effect of the change (on an interrupt controller for example) will not be seen until the end of the quantum when the SystemC scheduler is allowed to run. The delayed change will be unrealistic. TLM analysis port does not suffer this delay so is often used for this reason.

7 Tracing TLM activity

The OVPsim TLM2.0 interface can generate messages, controlled globally by setting environment variables, or locally by calling methods on the models. Set the environment variables to a non-null value or refer to the model header files (described in this document) for the method prototypes.

Model	Environment variable	Method	Trace
Processor	IMPERAS_TLM_CPU_TRACE	traceQuanta	start of each time slice
		traceBuses	each bus transaction
		traceBusErrors	each incomplete bus transaction
		traceSignals	signal value changes
Peripheral	IMPERAS_TLM_PSE_TRACE	traceBuses	each bus transaction
		traceBusErrors	each incomplete bus transaction
		traceSignals	signal value changes
MMC	IMPERAS_TLM_MMC_TRACE	traceMasters	each bus master transaction
		traceSlaves	each bus slave transaction

Note that when DMI (Direct Memory Interface) is enabled (which is by default on all OVP processors), there are no transactions to be traced.

Tracing of bus transactions is verbose and will reduce simulation performance.

8 What can go wrong

The following are a list of problems that can be encountered while building or running a TLM platform:

8.1 Spaces in filenames

nmake and other MSVC tools will not accept spaces in file-names.

Either install OVPsim, SystemC and TLM in a path without spaces (ie not "Program Files") or ensure that all paths are enclosed in double-quotes.

8.2 OVPsim version incompatibilities

Ensure that library, models and TLM2.0 interfaces are from the same version of OVPsim.

8.3 Environment problems

Check the values of environment variables:

- SYSTEMC
- TLM_HOME
- IMPERAS_VLNV,
- IMPERAS_HOME
- PATH

8.4 Compiling OSCI SystemC 2.2.0 with later versions of gcc

The OSCI SystemC simulator is available as source from <http://accellera.org/downloads/standards/systemc> and may be used with OVPsim, but when compiling the version 2.2.0 source with the latest version of GCC the following errors may be encountered:

```
../../../../src/sysc/utils/sc_utils_ids.cpp: In function 'int
sc_core::initialize()':
../../../../src/sysc/utils/sc_utils_ids.cpp:110: error:
\u2018getenv\u2019 is not a member of 'std'
../../../../src/sysc/utils/sc_utils_ids.cpp:111: error:
\u2018strcmp\u2019 was not declared in this scope
../../../../src/sysc/utils/sc_utils_ids.cpp: At global scope:
../../../../src/sysc/utils/sc_utils_ids.cpp:119: warning:
\u2018sc_core::forty_two\u2019 defined but not used
```

The solution is to add the following includes to the file
systemc-2.2.0/src/sysc/utils/sc_utils_ids.cpp:

```
#include "string.h"
#include "cstdlib"
```

Additionally, when compiling the OVP tlm modules you may see the following errors:

```
In file included from
/systemc-2.2.0/include/sysc/datatypes/bit/sc_lv.h:49,
from /systemc-2.2.0/include/sysc/communication/sc_signal_rv.h:61,
from /systemc-2.2.0/include/systemc:74,
from /TLM-2008-06-09/include/tlm/tlm.h:21,
from
/Imperas/ImperasLib/source/ovpworld.org/modelSupport/tlmMMC/1.0/tlm2.0/t
lmMmc.hpp:25,
from
/Imperas/ImperasLib/source/ovpworld.org/modelSupport/tlmMMC/1.0/tlm2.0/t
lmMmc.cpp:21:
/systemc-2.2.0/include/sysc/datatypes/bit/sc_lv_base.h: In member
function \u2018sc_dt::sc_logic_value_t sc_dt::sc_lv_base::get_bit(int)
const\u2019:
/systemc-2.2.0/include/sysc/datatypes/bit/sc_lv_base.h:310: error:
suggest parentheses around arithmetic in operand of \u2018| \u2019
/systemc-2.2.0/include/sysc/packages/boost/bind/placeholders.hpp: At
global scope:
/systemc-2.2.0/include/sysc/packages/boost/bind/placeholders.hpp:54:
error: \u2018<unnamed>::_1\u2019 defined but not used
/systemc-2.2.0/include/sysc/packages/boost/bind/placeholders.hpp:55:
error: \u2018<unnamed>::_2\u2019 defined but not used
...
```

The 'defined but not used' errors are due to an old version of Boost used in the SystemC source. To correct that problem the file
systemc-2.2.0/src/sysc/packages/boost/bind/placeholders.hpp may be edited
and line 28 changed from:

```
#if defined(__BORLANDC__)
```

to:

```
#if defined(__BORLANDC__) || defined(__GNUC__)
```

The warnings about parenthesis can be fixed by editing the make file to remove `-Werror` so that the warnings do not stop the compilation or by editing the indicated file and adding parenthesis where indicated.

```
##
```